



SPP '16 – MPI lab tasks



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Sergei Shudler



Overview – parallel sorting problem

- Lab purpose: Implement two parallel sorting algorithms and analyze the performance of the second one (the better one)
- Input:
 - N integers equally distributed among p processes
 - Process i has $n = N / p$ integers
- Output:
 - Permutation of the input numbers such that the local part in each process i is sorted
 - Each integer in process i is smaller or equal to all the other integers at a higher-ranked process j ($j > i$)

General remarks


- We provide header files and source templates
 - You need to implement specific functions
- Makefile and job script is provided
- Implementation has to be in C

Tasks

- Task 1 – Simple parallel sort
- Task 2 – Splitter-based parallel sort
- Task 3 – Complexity analysis
- Task 4 – Performance modeling

Simple parallel sort

- In Task 1 you should implement a simple parallel sort:
 - Each process sorts its local part of the array using std. C library
 - All p subarrays should be merged iteratively in parallel and the root should hold the sorted array in the end, for example ($p = 8$):



Simple parallel sort

- Results verification is provided
- Template: task1.c

Tasks



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Task 1 – Simple parallel sort
- Task 2 – Splitter-based parallel sort
- Task 3 – Complexity analysis
- Task 4 – Performance modeling

Splitter-based parallel sort - example

- Assume we have $p = 4$, $N = 32$, $n = 8$:

P0: 27 13 11 42 44 27 2 61

P1: 40 25 16 16 38 26 47 3

P2: 13 20 36 39 48 63 32 18

P3: 50 54 39 63 28 50 45 14

- Step 1 (sort locally):

P0: 2 11 13 27 27 42 44 61

P1: 3 16 16 25 26 38 40 47

P2: 13 18 20 32 36 39 48 63

P3: 14 28 39 45 50 50 54 63

Splitter-based parallel sort - example

- Step 2 (splitters – red boxes):

P0: 2 11 13 27 27 42 44 61

P1: 3 16 16 25 26 38 40 47

P2: 13 18 20 32 36 39 48 63

P3: 14 28 39 45 50 50 54 63

- Step 3 (new splitters – red boxes):

13 16 20 26 27 36 39 40 44 48 50 54

Splitter-based parallel sort - example

- Step 4 (splitters = 26, 39, 48):

P0: 2 11 13 27 27 42 44 61

P1: 3 16 16 25 26 38 40 47

P2: 13 18 20 32 36 39 48 63

P3: 14 28 39 45 50 50 54 63

Splitter-based parallel sort - example

- After the exchange:

P0: [2 11 13] [3 16 16 25] [13 18 20] [14]

P1: [26 38] [27 27] [32 36] [28]

P2: [39] [42 44] [40 47] [39 45]

P3: [50 50 54 63] [61] [48 63]

- Step 5 (local merge / sort):

P0: 2 3 11 13 13 14 16 16 18 20 25

P1: 26 27 27 28 32 36 38

P2: 39 39 40 42 44 45 47

P3: 48 50 50 54 61 63 63

Splitter-based parallel sort - remarks

- Details are in the exercise description
 - Use MPI_Gather, MPI_Bcast, MPI_Alltoall, and MPI_Alltoallv
 - Other collectives are possible, but not strictly necessary
- Implement an efficient results verification *verify_results_eff*
[Hint: Every process can send the biggest (local) element to the next neighbor]
- Template: task2.c

Tasks



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Task 1 – Simple parallel sort
- Task 2 – Splitter-based parallel sort
- Task 3 – Complexity analysis
- Task 4 – Performance modeling

- Analyze the complexity of the execution time and the memory of your implementation in Task 2
- Complexity of all the mandatory MPI collectives is provided:
 - MPI_Bcast: $O(m \cdot \log p)$, m – size of bcasted data
 - MPI_Gather: $O(\log p + m)$, m – the size of all the gathered data
 - MPI_Alltoall/v: $O(p + m)$, m – total amount of data sent from one process to all the others

Complexity analysis

- Simplifying assumption:
 - Every process has n integers
 - n is much bigger than p , so the terms with n will dominate over the terms in p
- Should not be too fine grained, but at the level of the algorithm steps provided in Task 2

Tasks



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Task 1 – Simple parallel sort
- Task 2 – Splitter-based parallel sort
- Task 3 – Complexity analysis
- Task 4 – Performance modeling

Performance modeling

- Purpose: use Extra-P to create a performance model for the runtime as a function of n
- You are provided with the job script *perf_analysis.sh*:
 - Runs the algorithm on 5 different values of n (on 32 processes)
 - Writes the output to *input.res*
- You have to run the script and run Extra-P on the *input.res*
- Output is produced in an *.xtrap* file, report the model you get in this file and compare to asymptotic complexity from Task 3